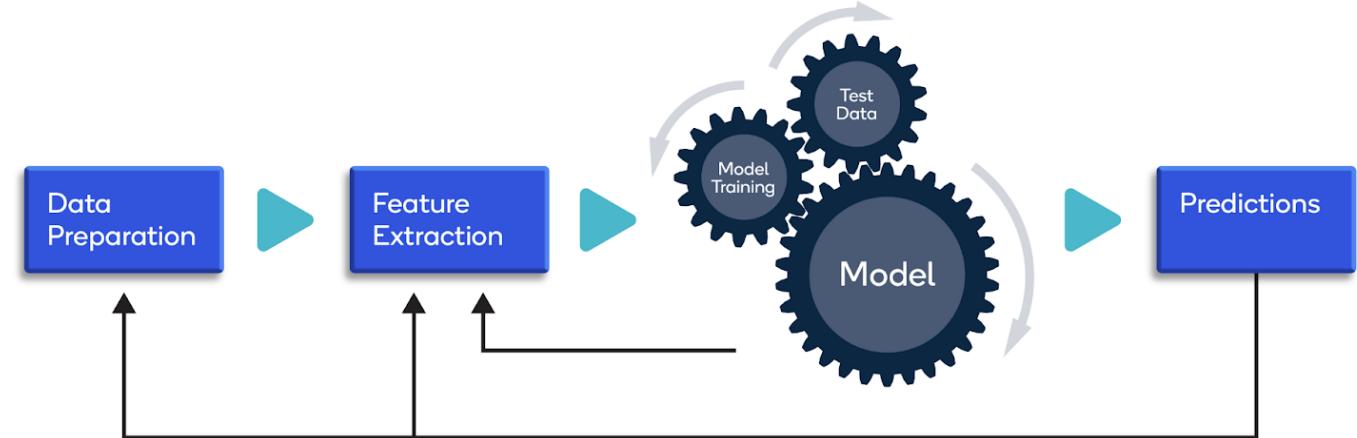


# CCS2213: Machine Learning

## *Topic 2*

## Machine Learning – Process & Metric



**Assoc. Prof. Dr Umi Kalsom Yusof**  
SCHOOL OF COMPUTING & INFORMATICS  
AlBukhary International University (AIU)

## Methodologies for Machine Learning



Fayyad's Knowledge Discovery in Databases (KDD)



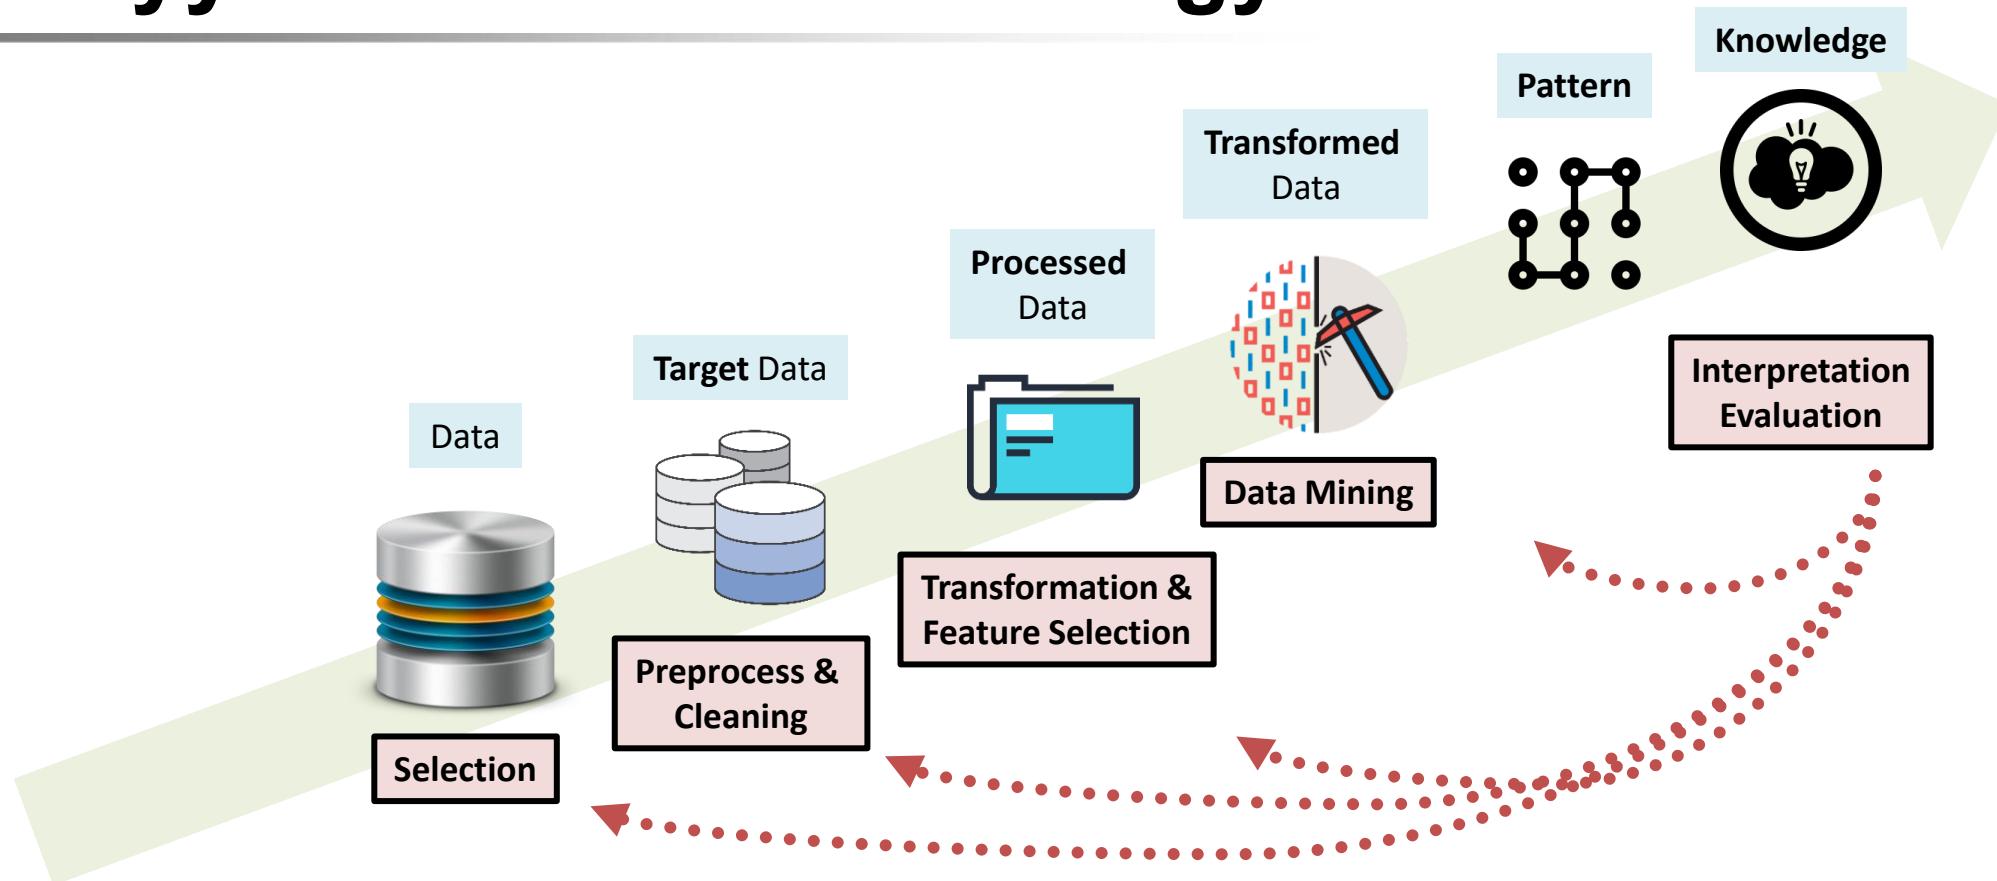
Sample, Explore, Modify,  
Model, Assess  
(SEMMA)

Data Mining Methodologies

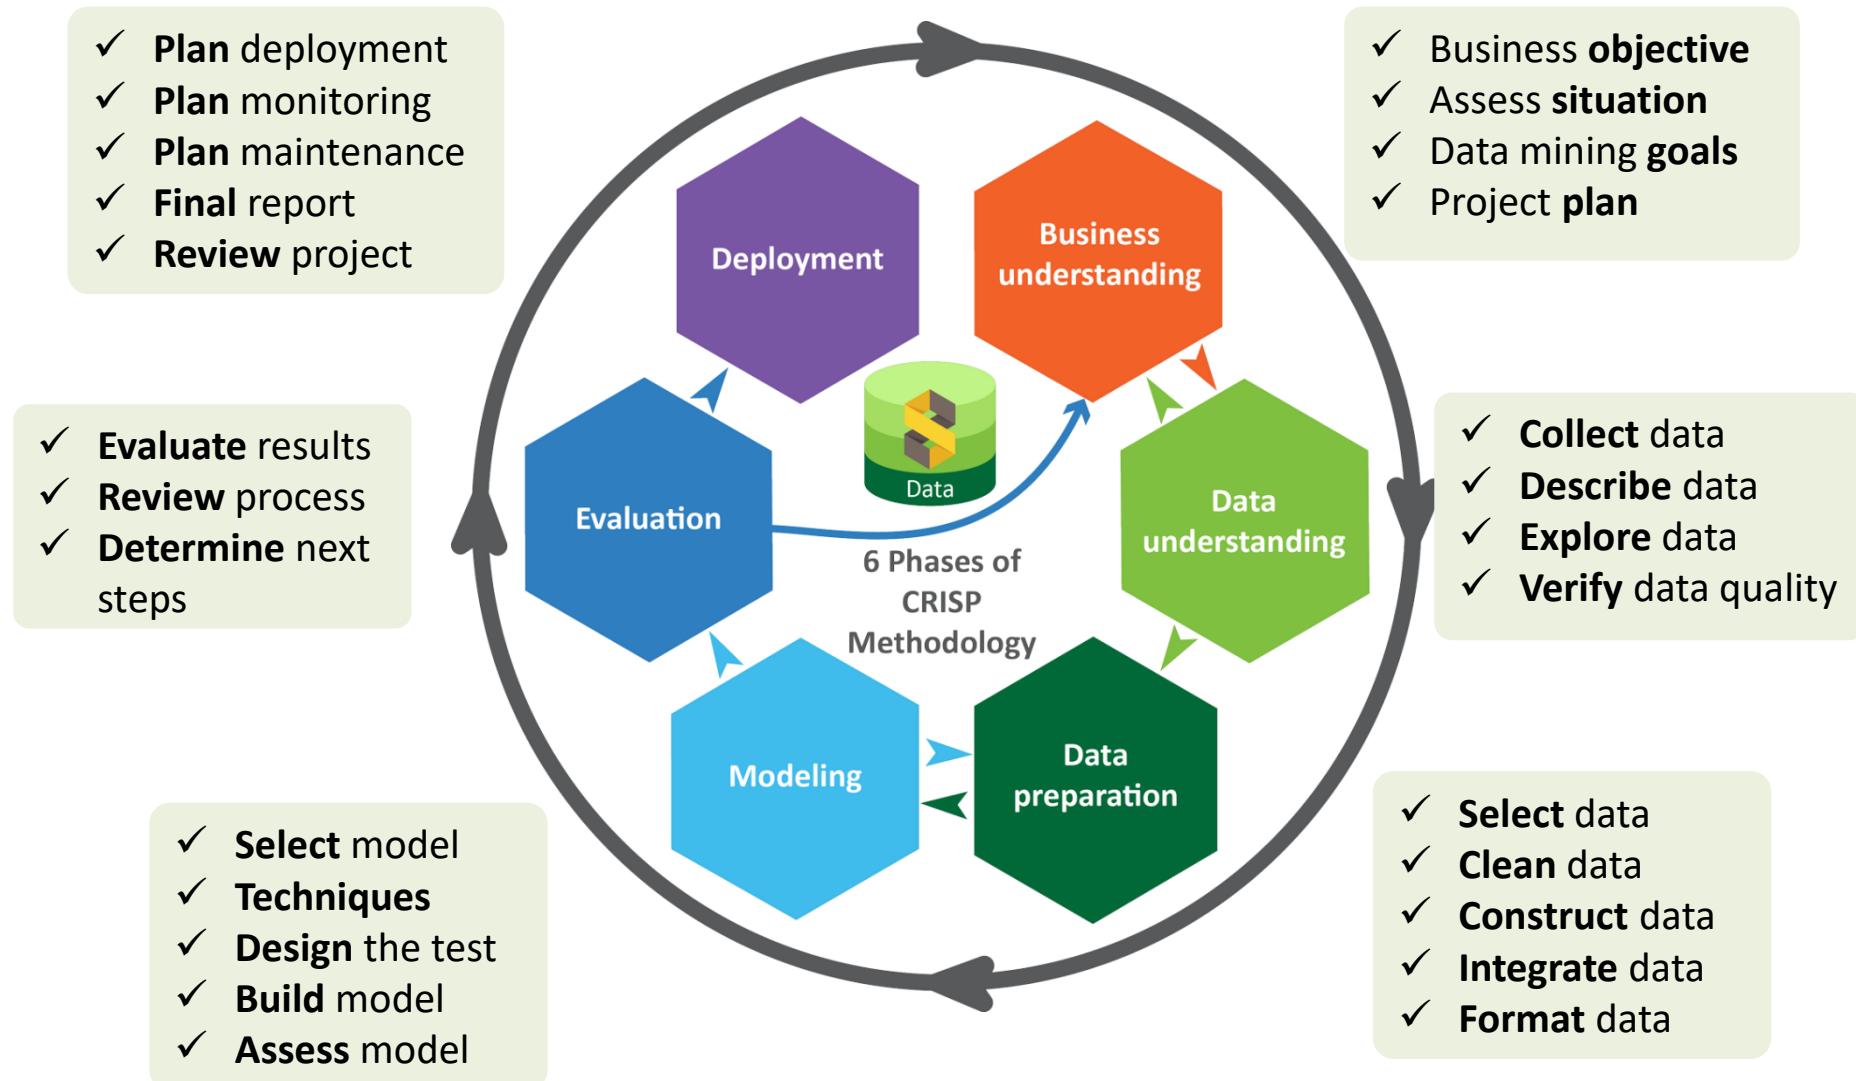


CRoss-Industry Standard  
Process for Data Mining  
(CRISP-DM)

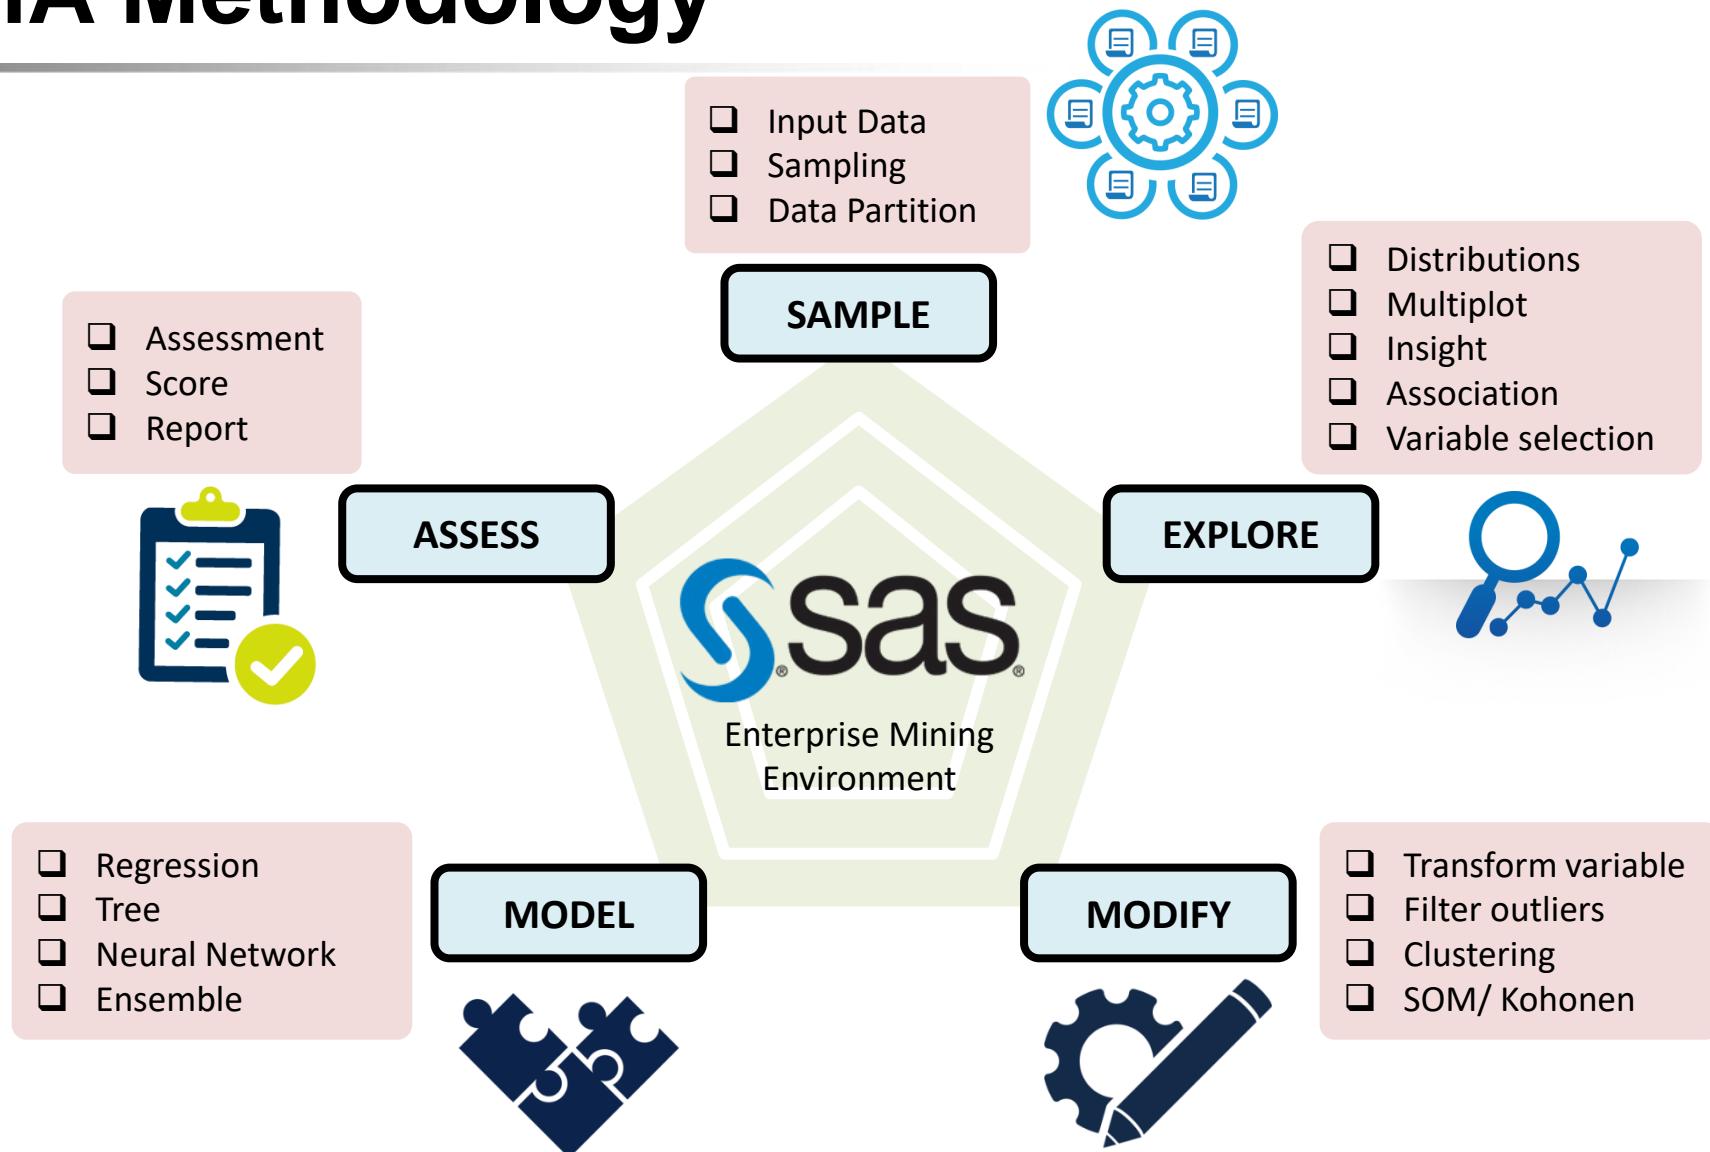
## Fayyad's KDD Methodology



## CRISP-DM Methodology

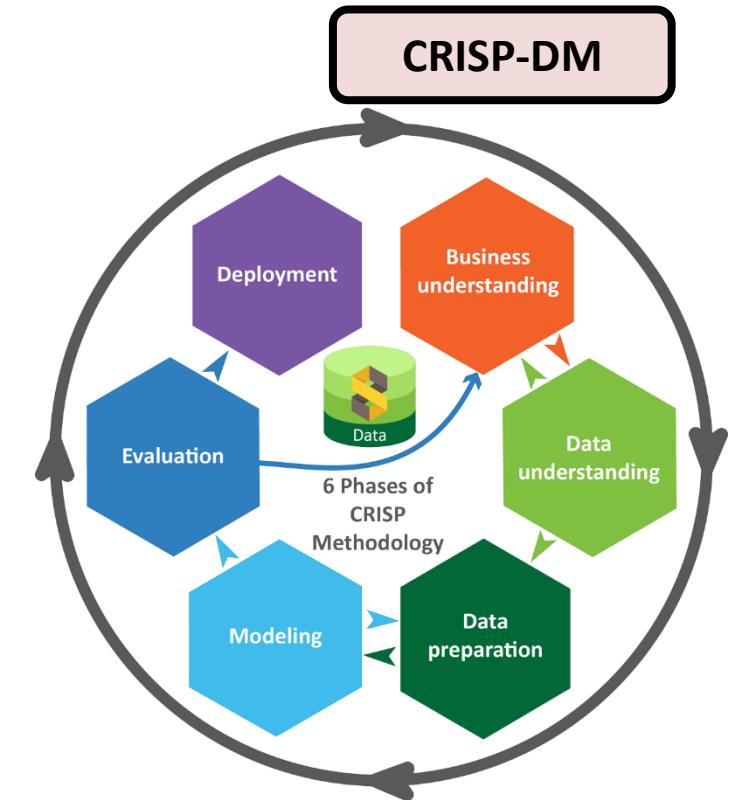
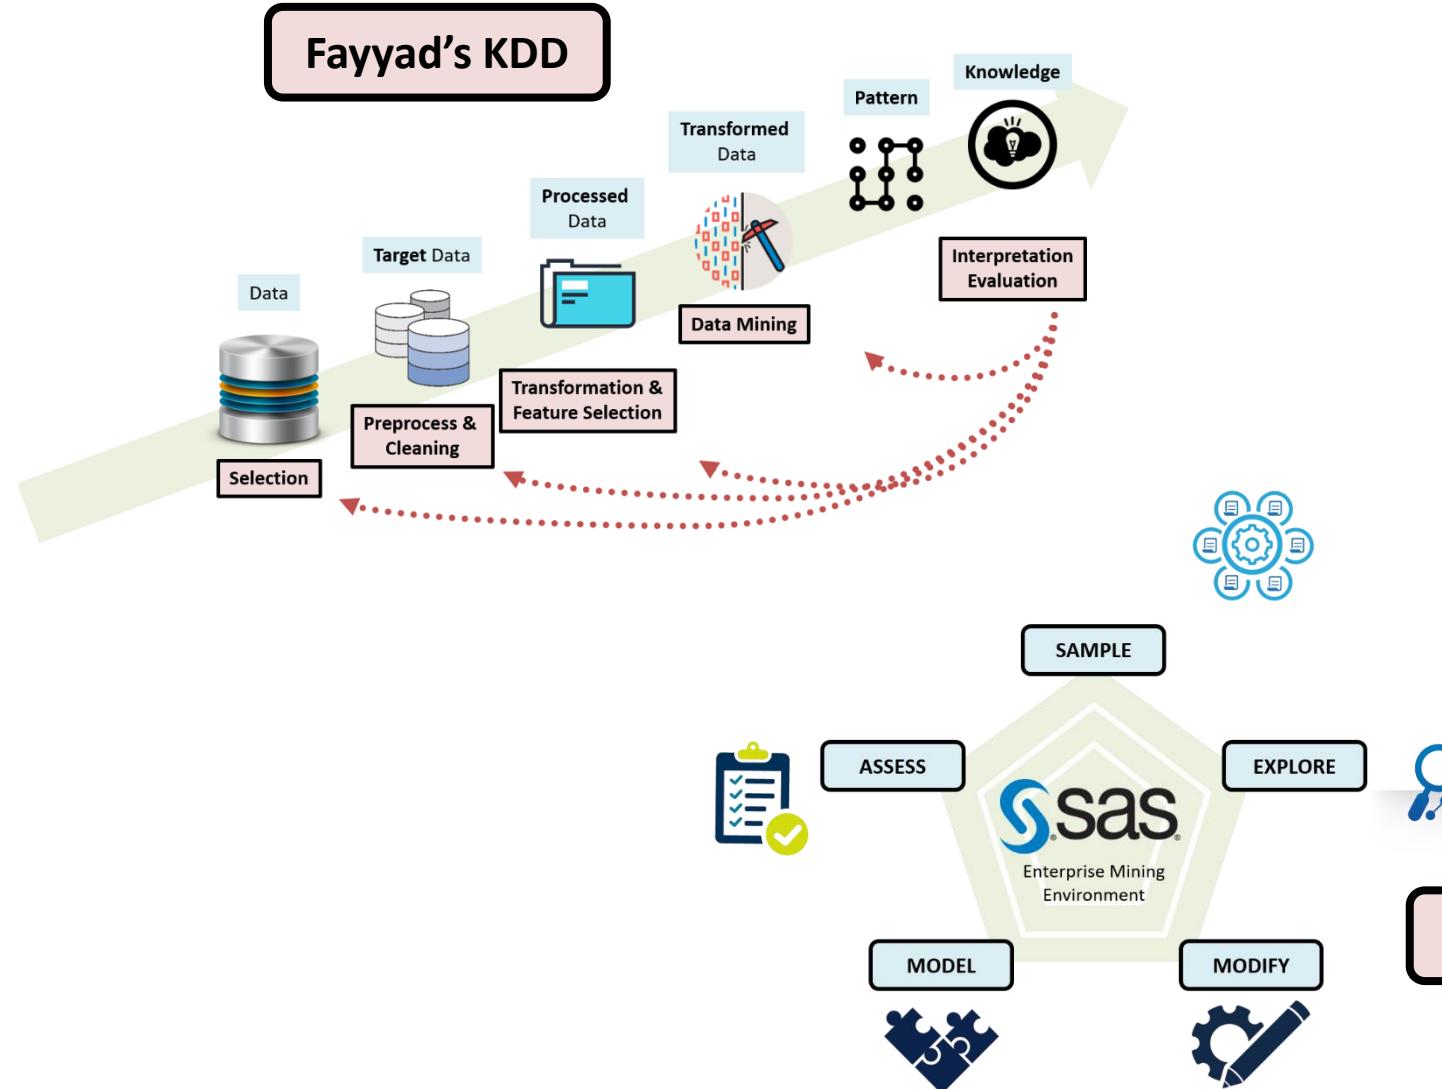


## SEMMA Methodology

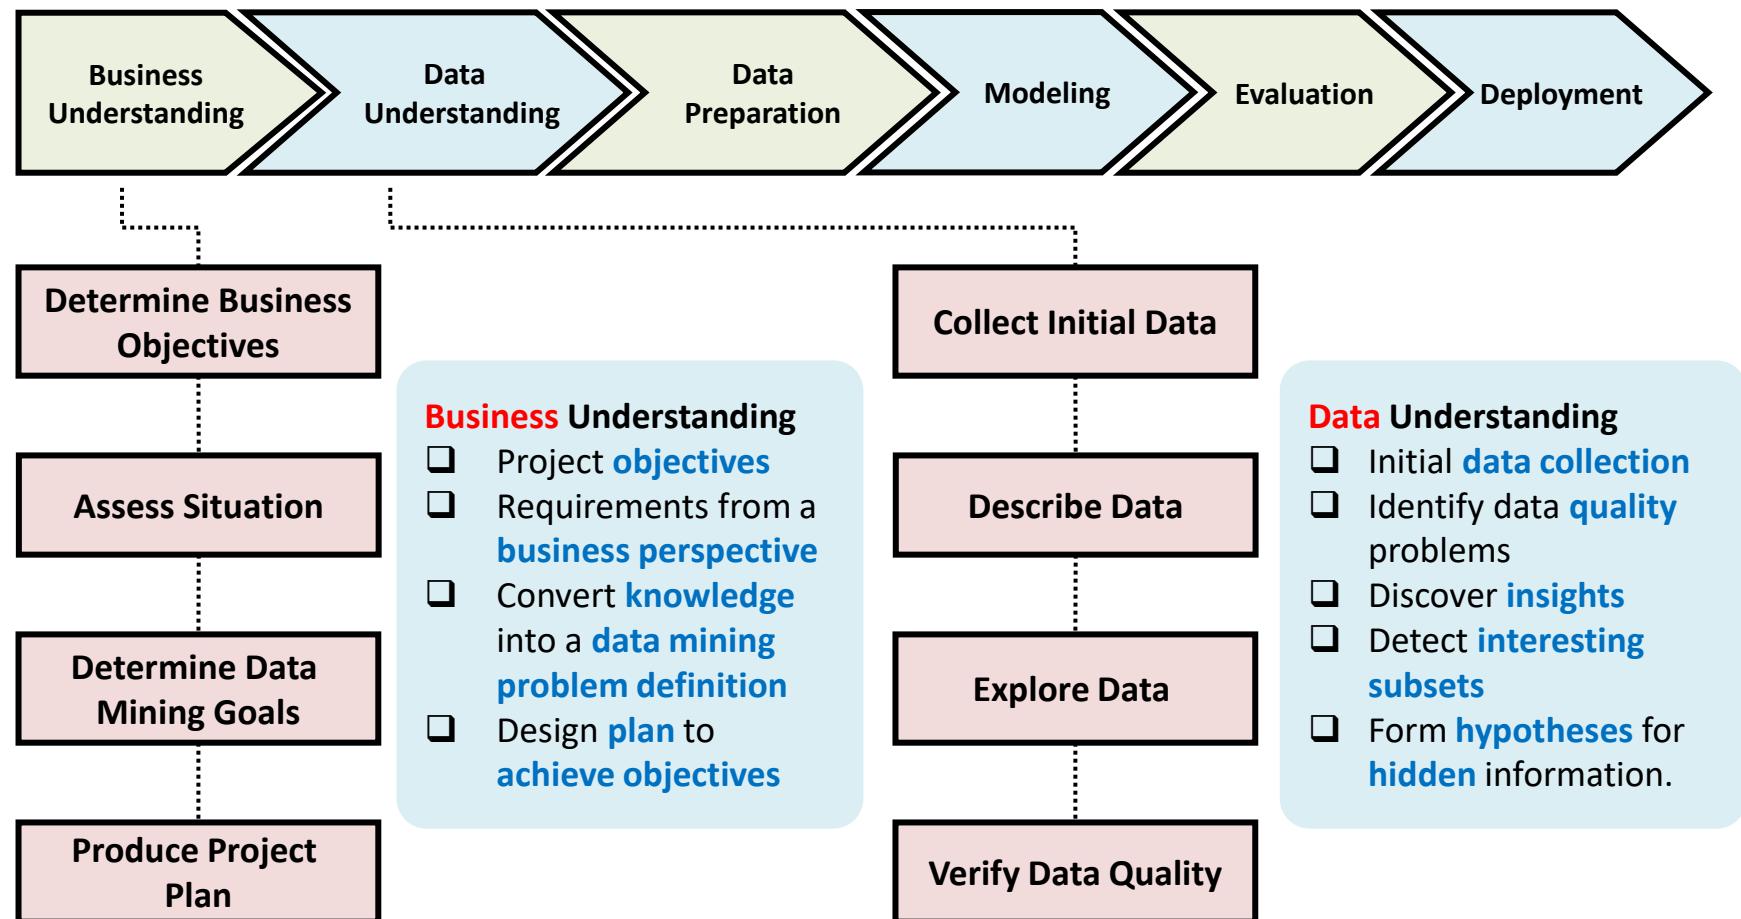


# Data Mining Process

## A Comparison



# Phases & Generic Tasks



# Motivation

# Model Performance Evaluation Techniques

Important to **evaluate classifier's generalization performance**:

Determine whether to employ the **classifier**

**Example:**

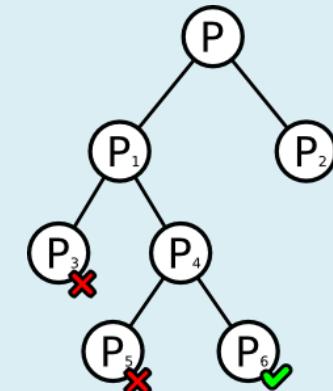
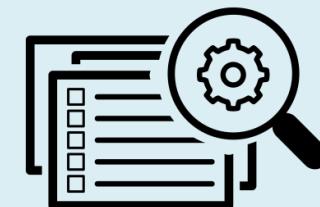
Learning the **effectiveness** of medical treatments from a **limited-size data**, it is important to **estimate** the **accuracy** of the **classifiers**

**Optimize** the classifier

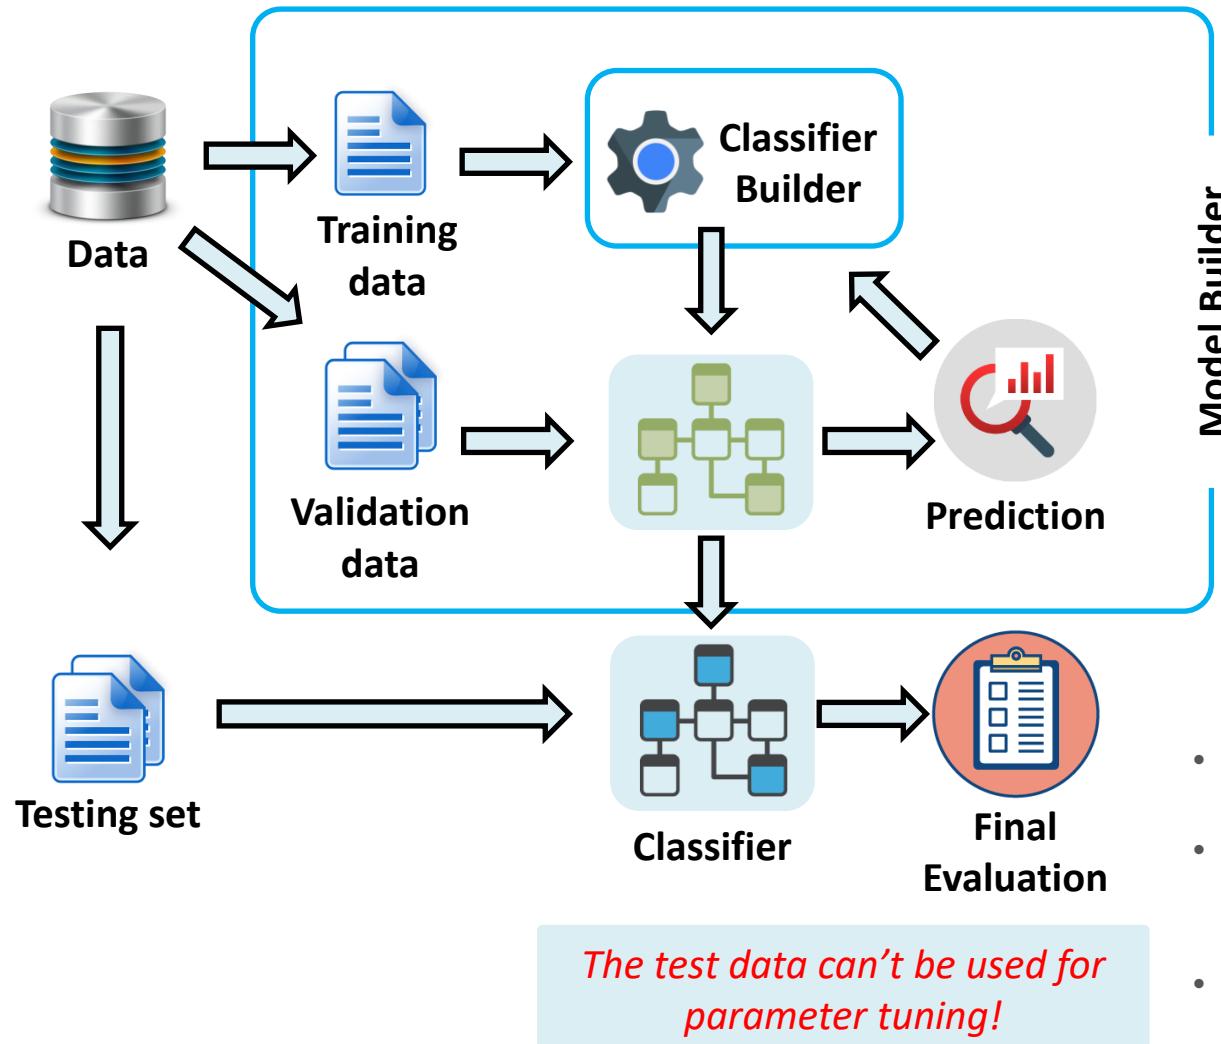
**Example:**

when **post-pruning** decision trees we must **evaluate** the **accuracy** of the decision trees on each **pruning step**

It is a ***data partitioning strategy*** so that you can effectively use your dataset to build a ***more generalized model***. The main intention of doing any kind of machine learning is to develop a more generalized model which can perform well on ***unseen data***.



# Classification: Train, Validation, Test Split



# Model Performance Evaluation Techniques

## Making The Most of The Data

Once evaluation is **complete**, all the **data** can be used to build the **final classifier**.

The **larger** the **training data** the **better** the classifier (but returns diminish).

The **larger** the **test data** the more **accurate** the error estimate.

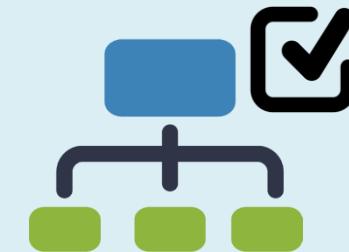
- **Training set:** A set of examples used for learning, that is to fit the parameters of the classifier.
- **Validation set:** A set of examples used to tune the parameters of a classifier, for example to choose the number of hidden units in a neural network.
- **Test set:** A set of examples used only to assess the performance of a fully-specified classifier.

# Model Performance Evaluation Techniques

For evaluating a **model's performance** and hyperparameter tuning

## Model evaluation techniques:

- **Training** data;
- **Independent** test data;
- **Hold-out** method;
- **k-fold** cross-validation method;
- **Leave-one-out** method;
- **Bootstrap** method;
- and many more...



# Model Performance Evaluation Techniques

## Training Data



The accuracy/metric estimates on the training data are not good indicators of performance on future data

New data will probably not exactly the same as the training data!

This measure the degree of classifier's overfitting (or underfitting).

## Independent Test Data



- Used when we have plenty of data
- Natural way of forming training and test data

### Example:

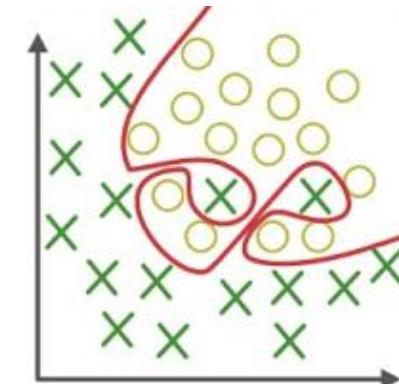
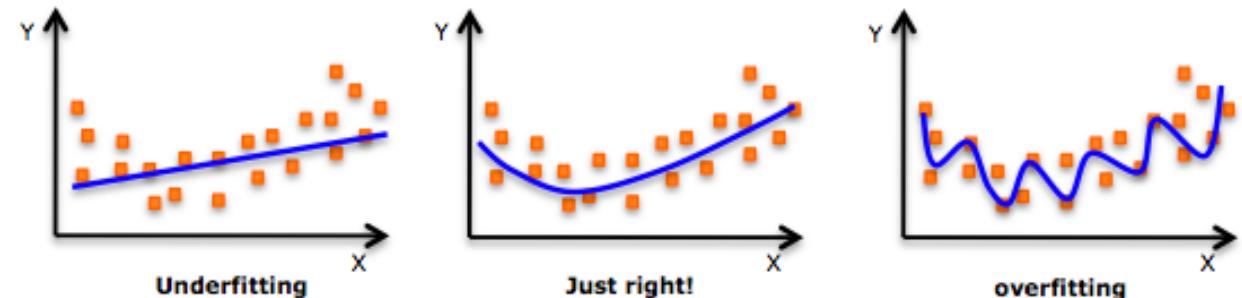
- Quinlan in 1987 reported experiments in a medical domain
- Trained on data from 1985
- Tested on data from 1986.

# What is Overfitting?

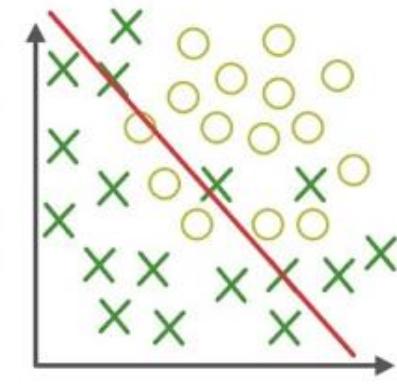
## Overfitting & Underfitting

Overfitting: A model that **models** the **training data** **too well**.

Underfitting: A model that can **neither** model the **training data** nor generalize to **new data**



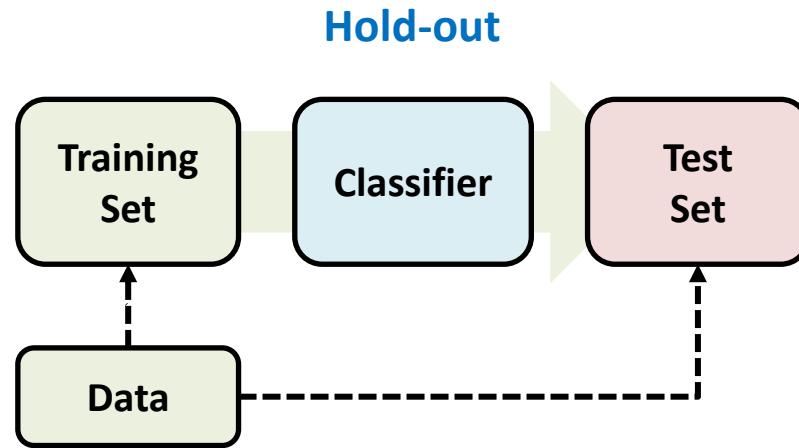
Over-fitting  
(forcefitting--too good to be true)



Under-fitting  
(too simple to explain the variance)

# Model Performance Evaluation Techniques

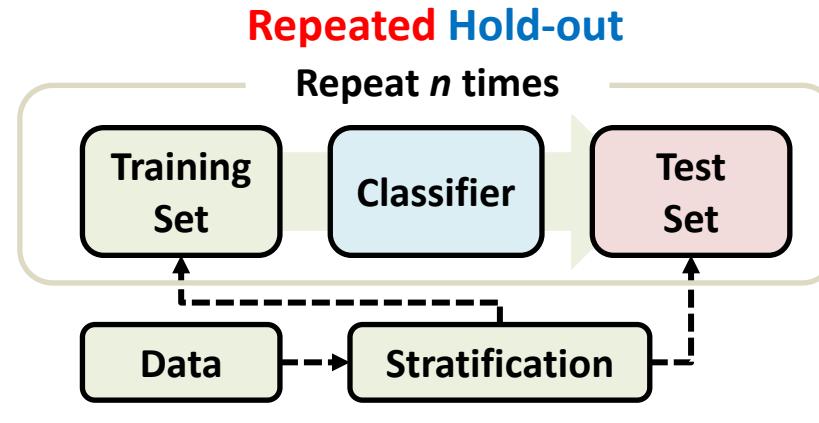
## Hold-out Method



The method **splits** the data into **training** data and **test** data (usually **2/3 for train, 1/3 for test**).

Then, **build** a classifier using the train data and **test** it using the test data.

Used when there are **thousands** of **instances** & several **hundred instances** from each **class**.



Holdout estimate can be made **more reliable** by **repeating** the process with **different subsamples**.

In each iteration, a certain **proportion** is **randomly** selected for **training** (possibly with **stratification**).

The **error rates** on the **different** iterations are **averaged** to yield an **overall error rate**.

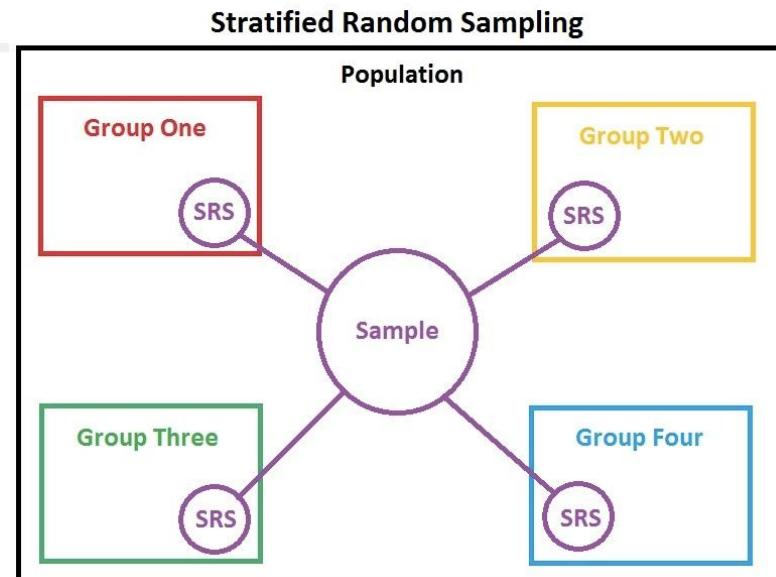
# Stratification

# Model Performance Evaluation Techniques

The **holdout** method **reserves** a certain amount for **testing** and uses the **remainder** for **training**.

For “unbalanced” datasets, samples **might not** be **representative** (**Few** instances or **none** of some classes).

**Stratified sample**  
An advanced version of **balancing** the data where **each class** is represented with **approximately equal proportions**



Stratification is the process of **dividing** members of the **population** into **homogeneous subgroups** before **sampling**.

The **strata** should be **mutually exclusive**

Every **element** in the population **must** be assigned to **only one stratum**.

The **strata** should be **collectively exhaustive**

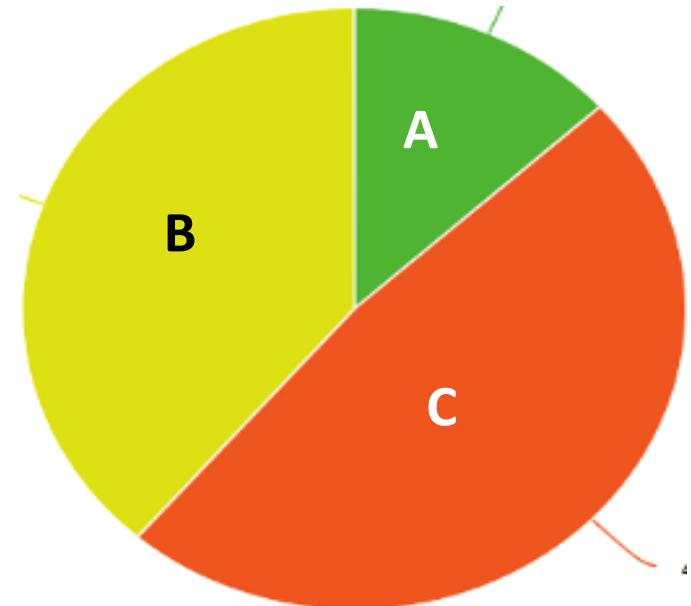
- No population **element** can be **excluded**.
- Simple random sampling is applied or;
- Systematic sampling is applied

This often **improves** the **representativeness** of the sample by **reducing** **sampling error**.

## Stratification

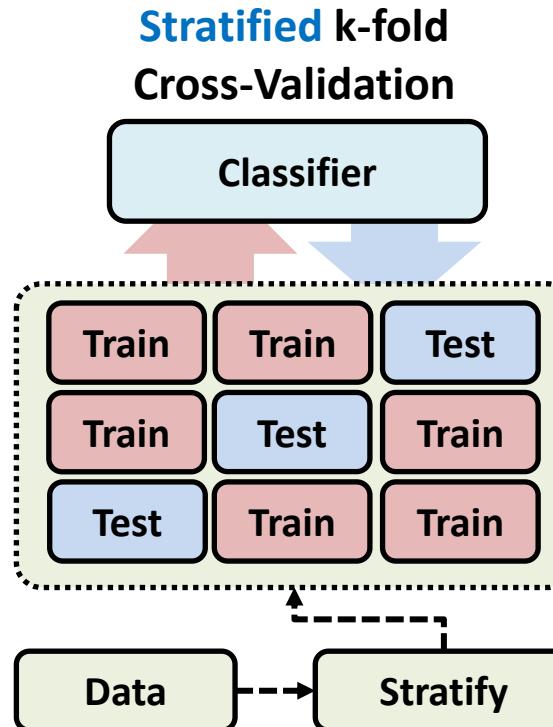
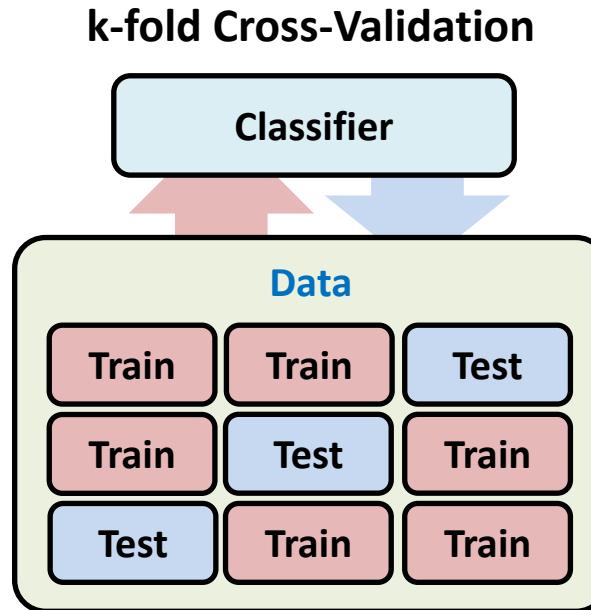
### Example

- Need to estimate **average number of votes** for **each candidate in an election**. The country has 3 towns:
  - Town A has **1 million factory workers**
  - Town B has **2 million office workers**
  - Town C has **3 million retirees**.
- A **random sample of size 60** over entire population will have some **chance of**:
  - The random sample **not well balanced** across these towns
  - Bias causing a **significant error** in estimation.
- **Random sample** of **10, 20 and 30** from Town A, B and C respectively
- Produce a smaller error in estimation for the same total size of sample.

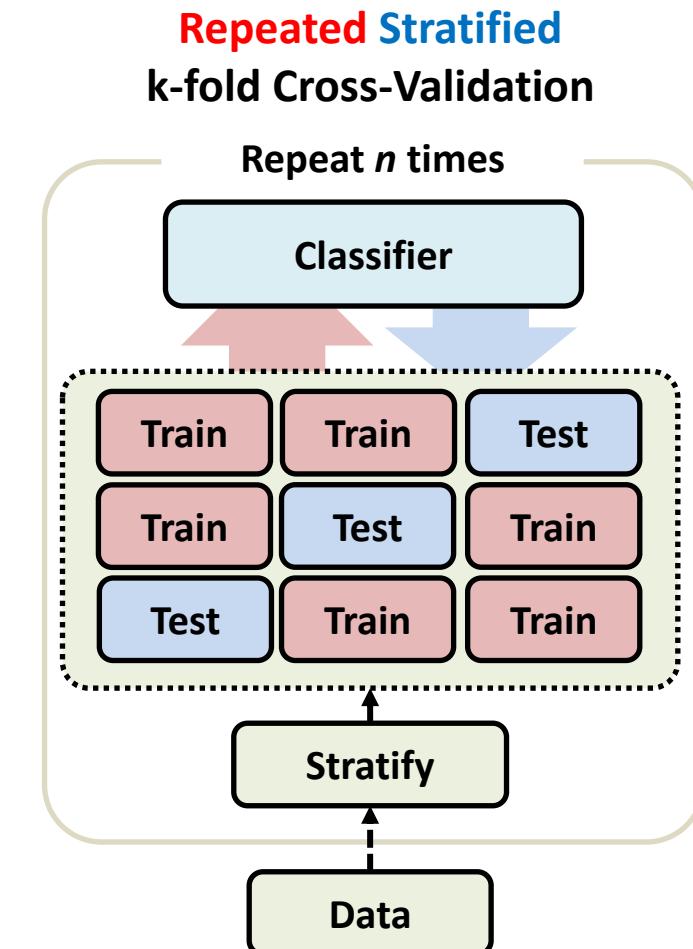


# Model Performance Evaluation Techniques

## k-Fold Cross-Validation



- Avoids overlapping test sets data is **equally split** into  $k$  subsets
- Some subset for **testing**, remainder for **training**.
- Recommended  $k = 10$  (best choice)
- Results are **averaged** (**reduces** the estimate's **variance**)
- Standard: **Stratified k-fold cross-validation**.
- Even better: **Repeated stratified k-fold cross-validation**.



## Leave-One-Out Cross-Validation

An **extreme** form of **cross-validation**

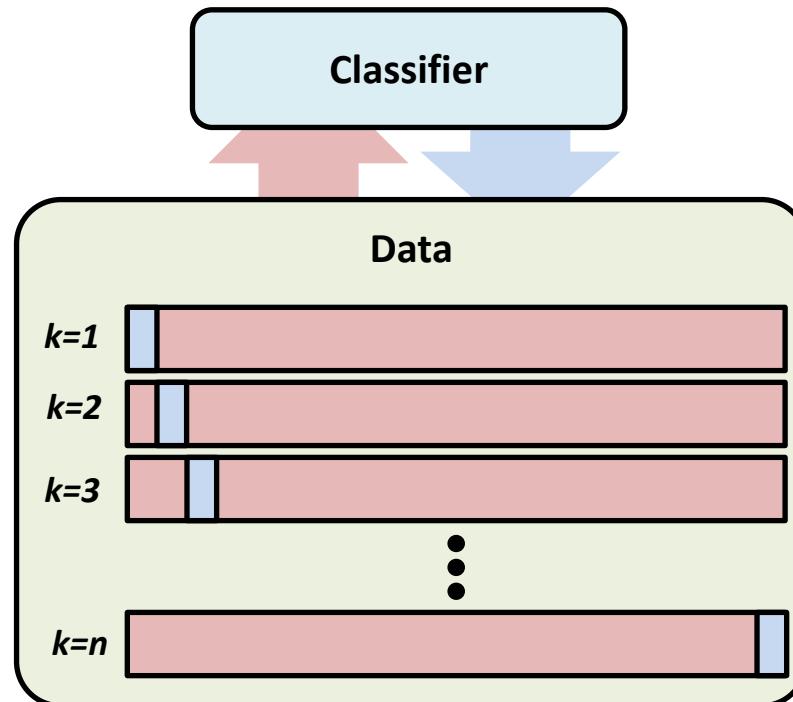
- Set **number of folds** to number of **training instances**;
- **N-1 training instances, 1 test instance**, and **build classifier N times**.

Make **best use** of the **data**

Involves **no** random sub-sampling

**Very computationally expensive**

Leave-One-Out Cross-Validation



# Classification measures

---

- Accuracy is only one measure ( $\text{error} = 1 - \text{accuracy}$ ).
- Accuracy is not suitable in some applications.
- In classification involving skewed or highly imbalanced data, e.g., network intrusion and financial fraud detections, we are interested only in the minority class.
  - High accuracy does not mean any intrusion is detected.
  - E.g., 1% intrusion. Achieve 99% accuracy by doing nothing.
- The class of interest is commonly called the positive class, and the rest negative classes.

## How to evaluate the Classifier's Generalization Performance?

- We test a **classifier** on some **test set**
- We derive at the end the following **confusion matrix**:

		<i>Predicted class</i>			
		Pos	Neg		
<i>Actual class</i>	Pos	<i>TP</i>	<i>FN</i>	<i>P</i>	<i>N</i>
	Neg	<i>FP</i>	<i>TN</i>		

# Classification Measures

- Accuracy =  $(TP+TN)/(P+N)$
- Error =  $(FP+FN)/(P+N)$
- Precision =  $TP/(TP+FP)$
- Recall/TP rate (Sensitivity) =  $TP/P$
- Specificity =  $TN/N$
- FP Rate =  $FP/N$

precision can be understood as the probability that a randomly chosen predicted positive instance would be relevant

recall is how close we are to a specific target on average.

- Accuracy =  $(100 + 50)/(120 + 80) = 150/200 = 75\%$
- Error =  $(30 + 20)/(120 + 80) = 50/200 = 25\%$
- Precision =  $TP/(TP+FP) = 100 / 130 = 76.9\%$
- Recall/TP rate =  $TP/P = 100/120 = 83.3\%$
- Specificity =  $TN/N = 50/80 = 62.5\%$
- FP Rate =  $FP/N = 30/80 = 37.5\%$

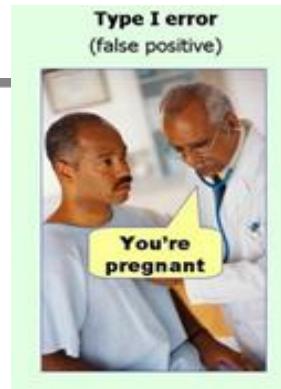
		<i>Predicted class</i>		<i>Recall</i>
		Pos	Neg	
<i>Actual class</i>	Pos	<i>TP</i>	<i>FN</i>	<i>P</i>
	Neg	<i>FP</i>	<i>TN</i>	<i>N</i>

**Precision**

		<i>Predicted class</i>		<i>Recall</i>
		Pos	Neg	
<i>Actual class</i>	Pos	<i>100</i>	<i>20</i>	<i>P</i>
	Neg	<i>30</i>	<i>50</i>	<i>N</i>

**Precision**

# Classification Measures



Accuracy:  $(TP + TN) / (P + N)$



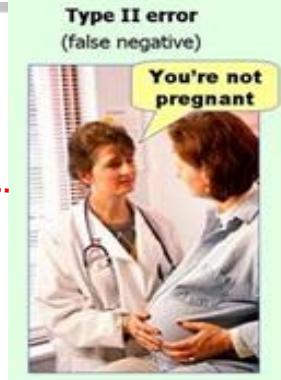
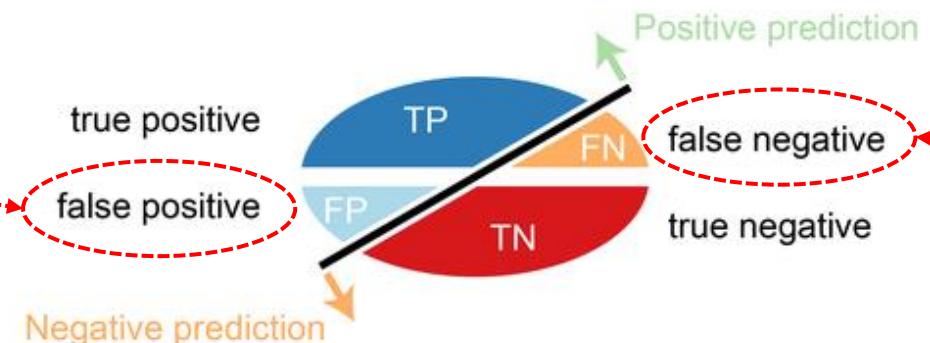
Accuracy is calculated as the total number of two correct predictions ( $TP + TN$ ) divided by the total number of a dataset ( $P + N$ ).

Sensitivity:  $TP / P$



Sensitivity is calculated as the number of correct positive predictions (TP) divided by the total number of positives (P).

## Four outcomes of a classifier



Error rate:  $(FP + FN) / (P + N)$



Error rate is calculated as the total number of two incorrect predictions ( $FN + FP$ ) divided by the total number of a dataset ( $P + N$ ).

Precision:  $TP / (TP + FP)$



Precision is calculated as the number of correct positive predictions (TP) divided by the total number of positive predictions ( $TP + FP$ ).

False positive rate:  $FP / N$



False positive rate is calculated as the number of incorrect positive predictions (FP) divided by the total number of negatives (N).

## $F_1$ -value (also called $F_1$ -score)

- It is **hard** to compare two classifiers using two measures.  $F_1$  score combines **precision** and **recall** into one measure

Precision =  $TP/(TP+FP)$

Recall/TP rate (**Sensitivity**) =  $TP/P$

$$F_1 = \frac{2pr}{p+r}$$

$F_1$ -score is the harmonic mean of precision and recall.

$$F_1 = \frac{2}{\frac{1}{p} + \frac{1}{r}}$$

- The **harmonic** mean of two numbers tends to be closer to the smaller of the two.
- For  $F_1$ -value to be **large**, both  $p$  and  $r$  must be **large**.

# Cost Sensitive Learning

## Cost Matrices & Classification

**Goal** of machine learning is based around making a computer **generalize** its **observation**.

The **measure of performance** of a machine learning algorithm is based on its **accuracy** of **classifying** a data set.

In practice, **different types of classification errors** often incur **different costs**

### Examples:

- Terrorist **profiling**
- “Not a terrorist” correct 99.99% of the time
- Loan **decisions**
- Fault **diagnosis**
- **Promotional** mailing

		Cost Matrices	
		Predicted class	
		Positive	Negative
True class	Positive	TP Cost	FN Cost
	Negative	FP Cost	TN Cost

\*Usually, **TP Cost** and **TN Cost** are set equal to **0**!

If the classifier **outputs** class probability, adjust it to **minimize** the **expected prediction cost**.

**Expected cost** is computed as **dot product** of class probabilities **vector** with appropriate **column** in cost matrix.

# Cost Sensitive Learning

## Cost Matrices & Classification

### Example

Assume that a classifier returns for an instance probs  $p_{pos} = 0.6$  and  $p_{neg} = 0.4$ .

The expected cost if the instance is classified as positive:

$$0.6 * 0 + 0.4 * 10 = 4$$

The expected cost if the instance is classified as negative:

$$0.6 * 5 + 0.4 * 0 = 3$$

- Simple methods for cost sensitive learning
- Resampling of instances according to costs;
  - Weighting of instances according to costs.

		Predicted class	
		Positive	Negative
True class	Positive	0	5
	Negative	10	0
		Predicted class	
		Positive	Negative
True class	Positive	$0.6 * 0$	$0.6 * 5$
	Negative	$0.4 * 10$	$0.4 * 0$

# ROC Curve

How do we pick the probability threshold that gives us the **best performance** for the situation that we want?

ROC Curve: Receiver operating characteristic curve

It is a graphical representation of how two of these metrics (**Sensitivity** or **Recall** and the **Specificity**) vary as we change this probability threshold.

- Recall/TP rate (**Sensitivity**) = TP/P
- **Specificity** = TN/N

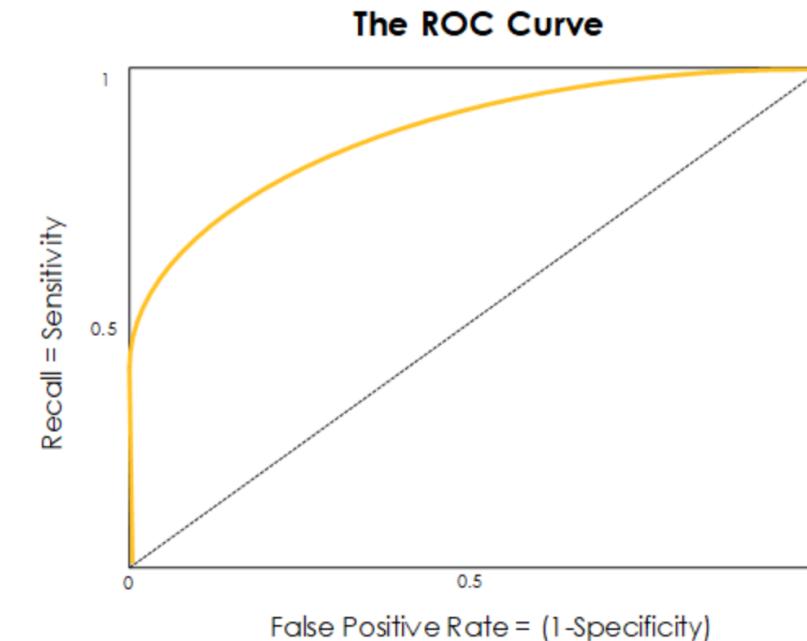
		<i>Predicted class</i>	
		Pos	Neg
<i>Actual class</i>	Pos	<i>TP</i>	<i>FN</i>
	Neg	<i>FP</i>	<i>TN</i>

- It is a plot of the **true positive rate (TPR)** against the **false positive rate (FPR)**.
- **True positive rate:**

$$TPR = \frac{TP}{TP + FN}$$

- **False positive rate:**

$$FPR = \frac{FP}{TN + FP}$$



# ROC Curves and Analysis

A ROC (Receiver Operating Characteristics) curve :

Classifier 1

True	Predicted	
	pos	neg
pos	40	60
neg	30	70

Classifier 2

True	Predicted	
	pos	neg
pos	70	30
neg	50	50

Classifier 3

True	Predicted	
	pos	neg
pos	60	40
neg	20	80

Classifier 1  
TPr = 0.4  
FPr = 0.3

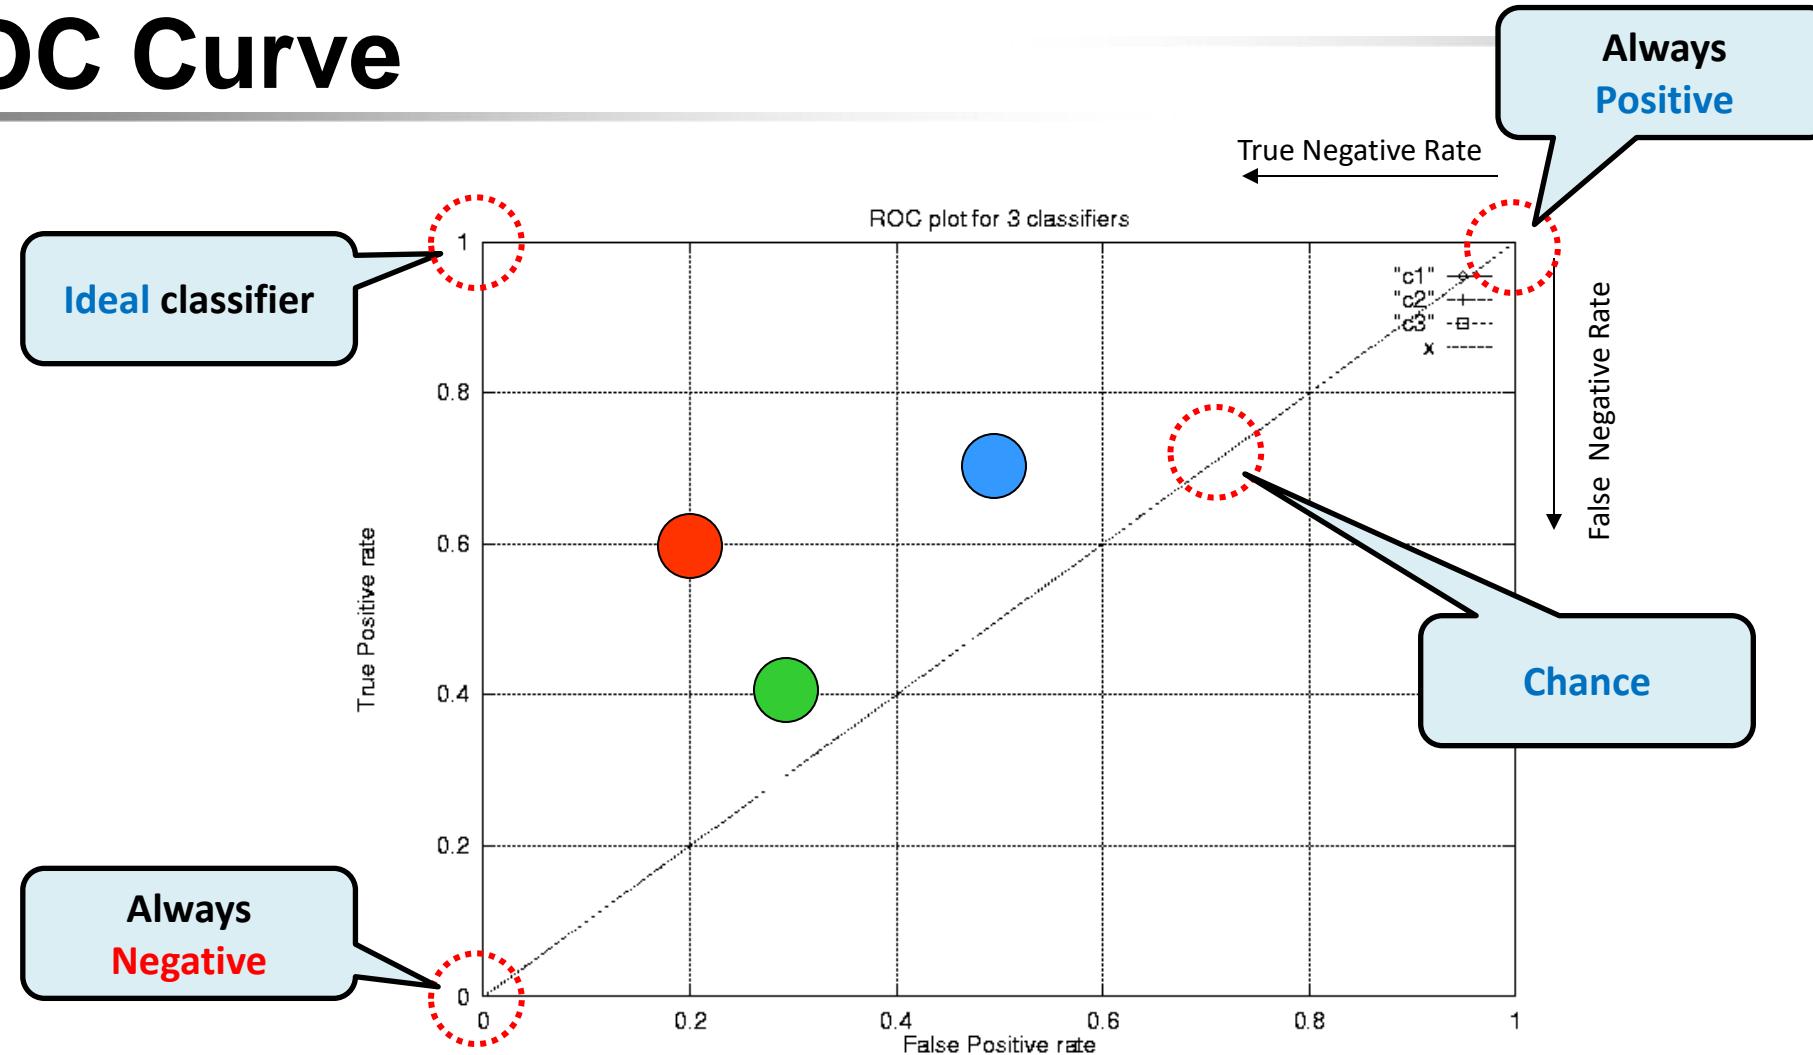
Classifier 2  
TPr = 0.7  
FPr = 0.5

Classifier 3  
TPr = 0.6  
FPr = 0.2

TPr – True Positive rate

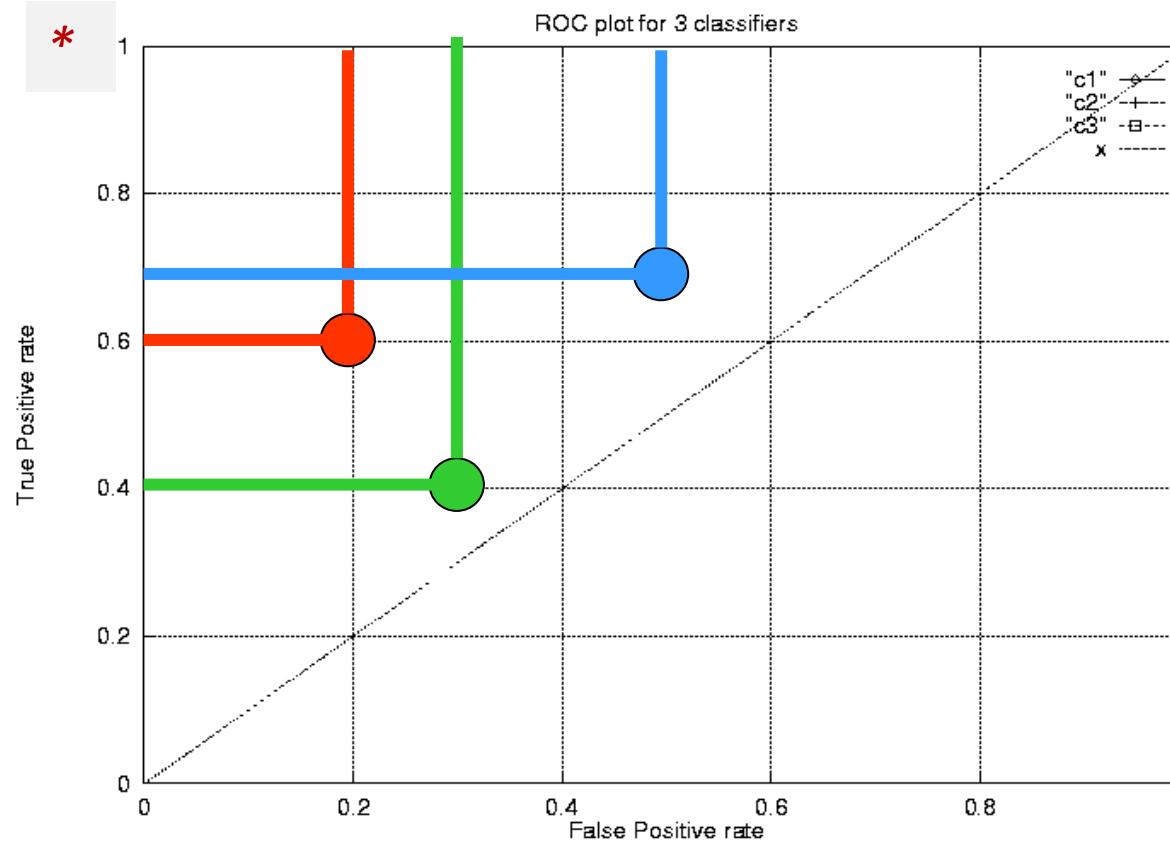
FPr – False Positive rate

# ROC Curve



# Dominance in the ROC Curve

Which one is the best option?



*Classifier A dominates classifier B if and only if  $TPr_A > TPr_B$  and  $FPr_A < FPr_B$ .*

*How lose each classifier to point \*?  
How to calculate?*

*Can also use Euclidean distance formula to calculate how far from (0, 1)*

---

---

**Thank you**